Reinforcement Learning, Tutorial 06

Philipp Kratzer

Machine Learning and Robotics Lab

University of Stuttgart
Germany

June 17th, 2020

Outline

Announcements

- Next exercise sheet is available
- In total we plan to have 10 exercise sheets ($\sim$100 points)
- You need at least 50% to pass the exercises $\rightarrow$ $\sim$50 points

Outline

1a

Do you think a n-step bootstrapping method could do as well as the Dyna-Q method in figure below?

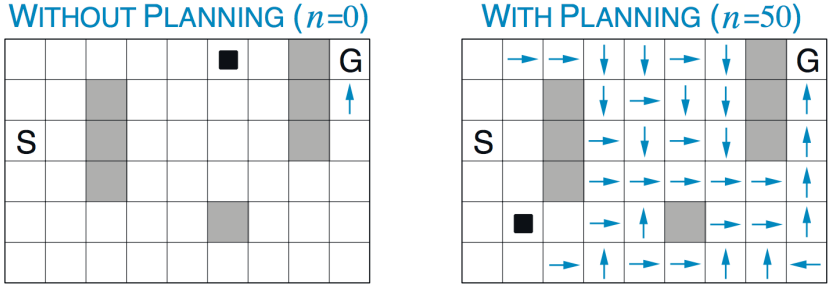

1b

Why did Dyna-Q+ perform better in both phases?

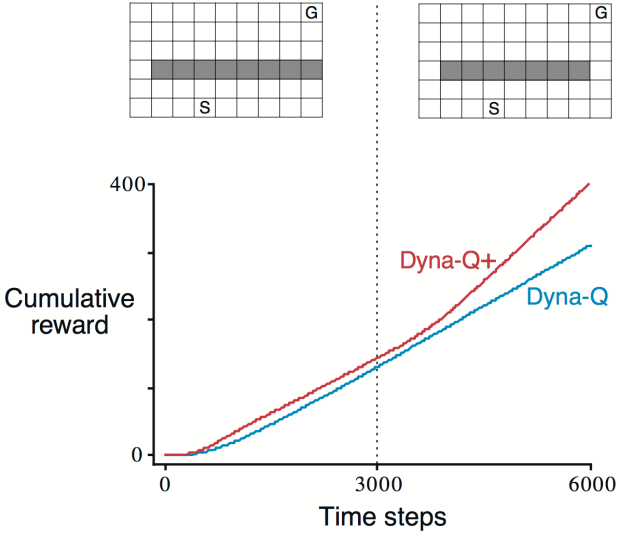

2

Implement n-step Sarsa and plot the performance over α for different choices of n.

Outline

Outlook

- Next exercise sheet available
- Programming will be on the MountainCar environment